Assessing the Limits of the Distributional Hypothesis in Semantic Spaces

Trait-based Relational Knowledge and the Impact of Co-occurrences

Mark Anderson and Jose Camacho Collados

Research Questions

General

What is required in data for models to capture meaningful representations of natural language?

Specific

What is the impact of co-occurrences of concepts and their traits on the ability of semantic spaces to encode this type of relational knowledge?

Trait-based Relational Knowledge

Traits/attributes

- Traits commonly associated with concepts.
- Split into types (colour, shape&size, material, components, tactile)
- E.g. The ubiquitous yellow banana.

Datasets

- McRae[1]
- Norms[2]
- McRae (Spanish)

---

[1] Ken McRae, George S. Cree, Mark S. Seidenberg, and Chris McNorgan. 2005. *Semantic feature production norms for a large set of living and nonliving things.*Behavior Research Methods, 37:547–559.

[2]Barry Devereux, Lorraine K. Tyler, Jeroen Geertzen,and Billi Randall. 2014. *The centre for speech, language and the brain (CSLB) concept property norms.* Behavior Research Methods, 46:1119 – 1127

Example table — McRae

	trait type	N_C		N_T	
	colour	148		7	green (32), brown (32), black (24), white (21), red (16), yellow (13), orange (10)
	components	110		6	handle (39), legs (19), wheels (14), leaves (14), seeds (13), doors (11)
	materials	144		4	metal (79), wood (43), cotton (11), leather (11)
	size & shape	234		4	small (83), large (70), long (44), round (37)
	tactile	117		7	heavy (21), soft (19), furry (18), sharp (17), hard (16), juicy (16), slimy (10)

Co-occurences

- Sentence — e.g. *The **bananas** were ubiquitous in that town, and **yellow** jackets were the outerwear of choice for the inhabitants.*

- Window (±5 tokens of concept) — e.g. *He wore a [**yellow** t-shirt while eating a **banana** in the garage of a] kind stranger.*

- Syntactic — e.g. *If **bananas** aren't **yellow**, nobody will be interested in eating them, except for perhaps some questionable folk.*

Experimental Details 1

Corpora

- **UMBC** 3.4B Tokens (English)
- **Wikipedia** Dump 2.5B/0.6B Tokens (English/Spanish)
- **ES1B** 1.4B Tokens (Spanish)

Removing Co-occurrences

- Use 80% of corpus as training data nd keep 20% as reserves.
- Sentences with co-occurrences removed from training data.
- Replaced with random sentence from reserves.

Experimental Details 2

Vector Space

- **CBOW Word2Vec**
- Faster than skip-gram.
- Trained models with and without co-occurrences.
- For each separate trait type (e.g. colour)

Classifiers

- Multi-class (predict **yellow** given **banana** as input, as embedding from model).
- Binary-class (predict **related** or **not related** with $e_{concept}$-$e_{relation.}$ as input).
- Used SVMs.

Results?

McRae (English) — Multi-class

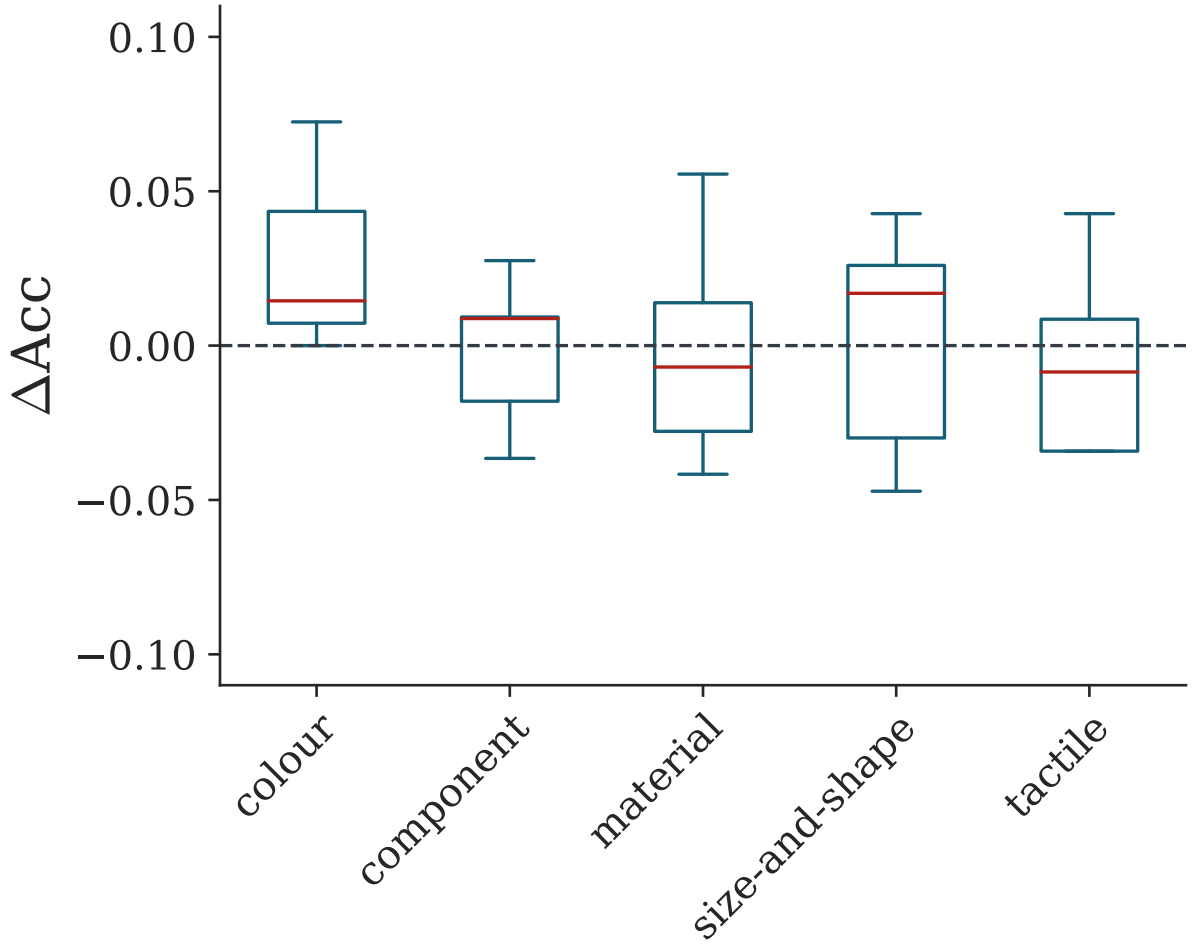

Figure 1: Distribution of multi-class accuracy differences between models trained with and without co-occurrences.

Would anticipate mean difference to be > 0, if co-occurrences help. Only true for colour trait type.

Norms (English) — Multi-class

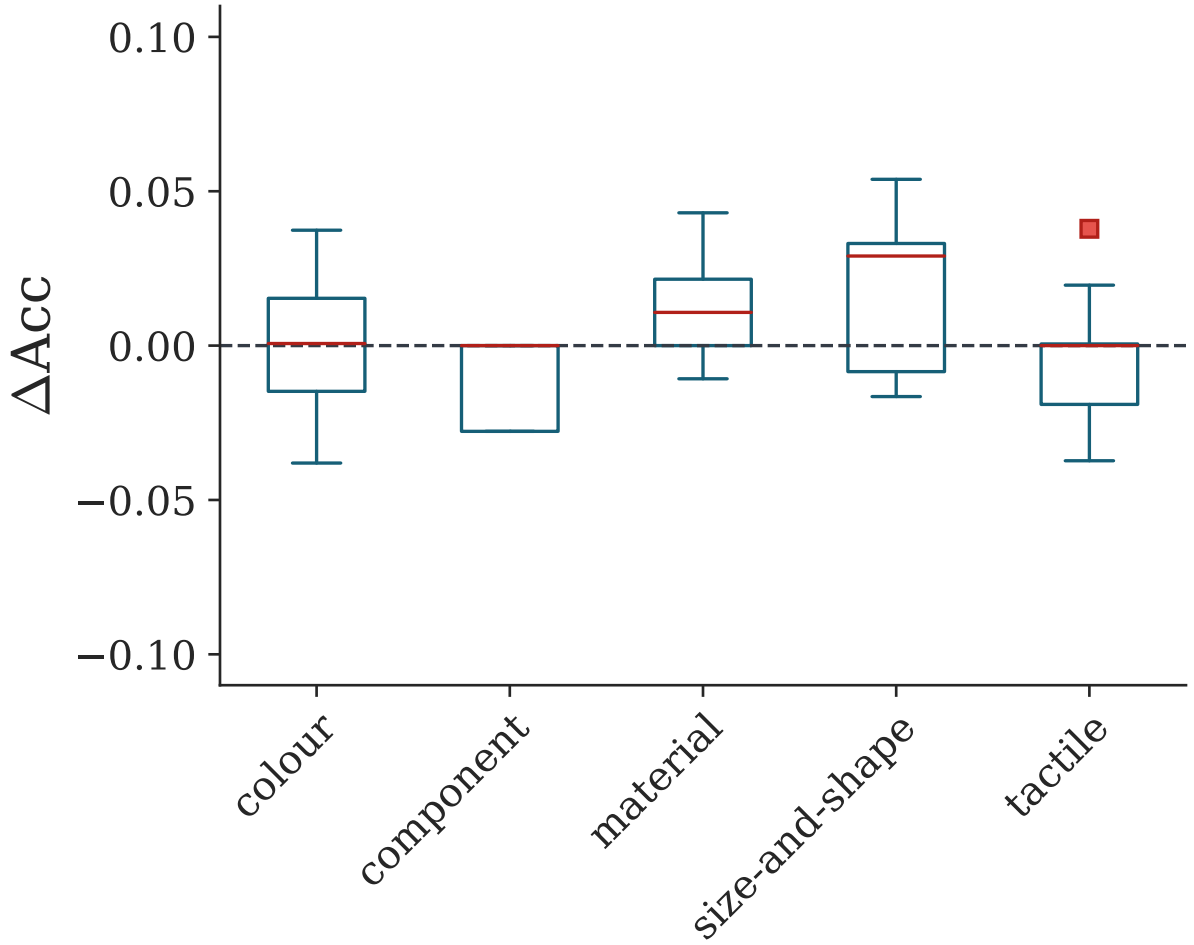

Figure 2: Distribution of multi-class accuracy differences between models trained with and without co-occurrences.

McRae (Spanish) — Multi-class

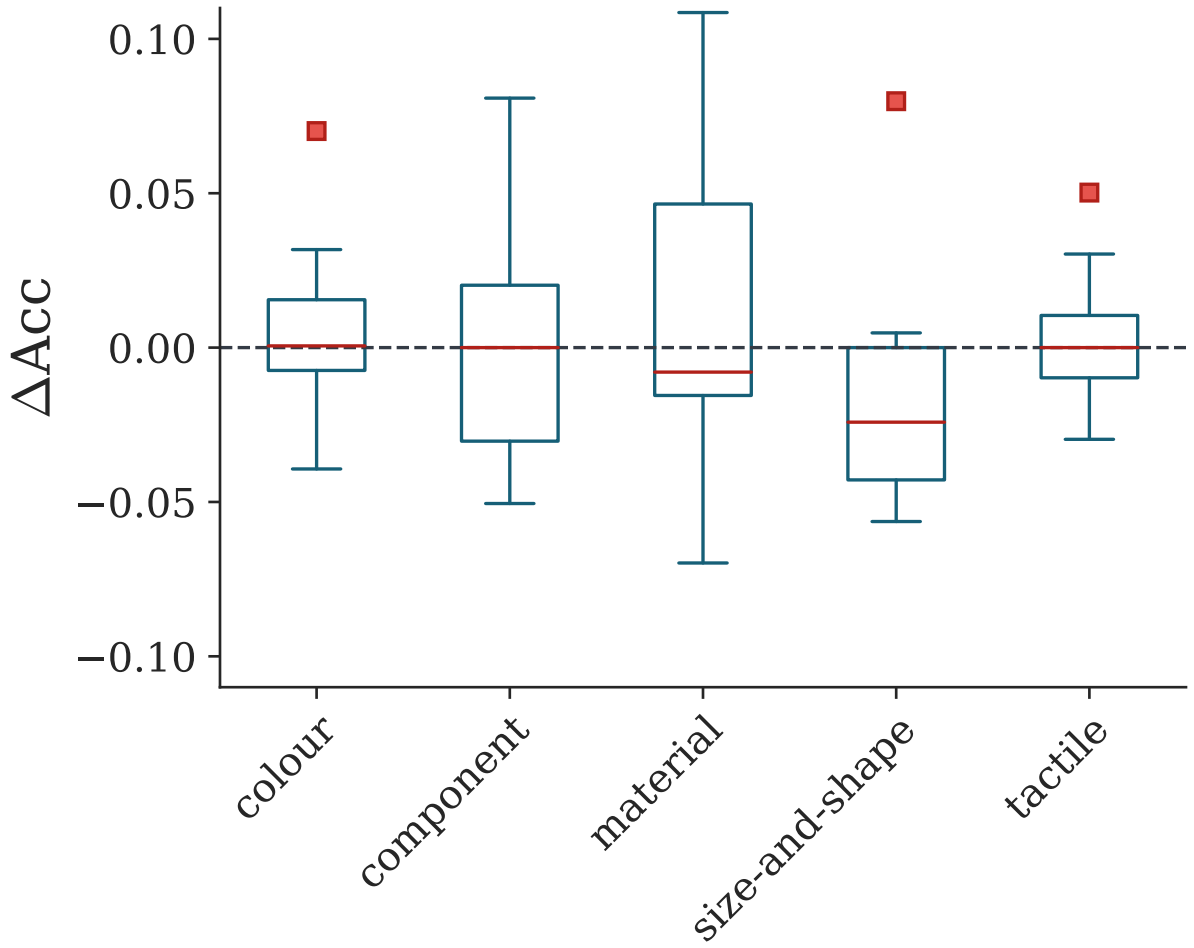

Figure 3: Distribution of multi-class accuracy differences between models trained with and without co-occurrences.

McRae (English) — Multi-class

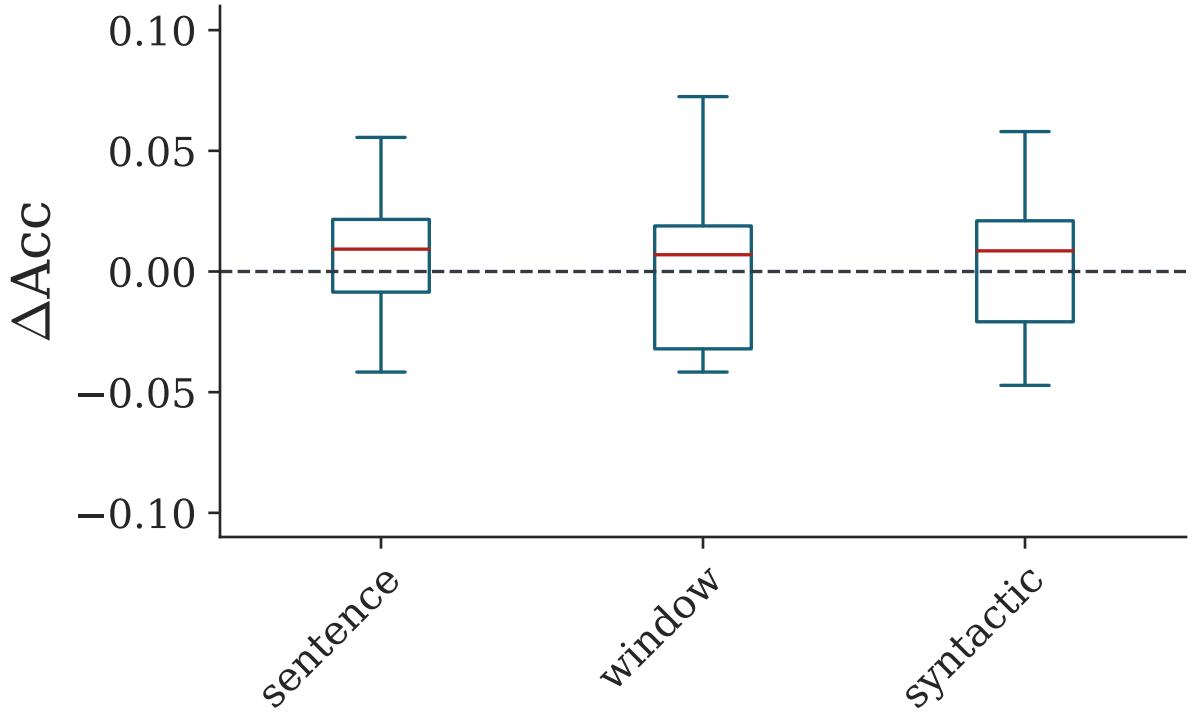

Figure 4: Distribution of multi-class accuracy differences between models trained with and without co-occurrences.

Mean differences around zero for all methods.

Norms (English) — Multi-class

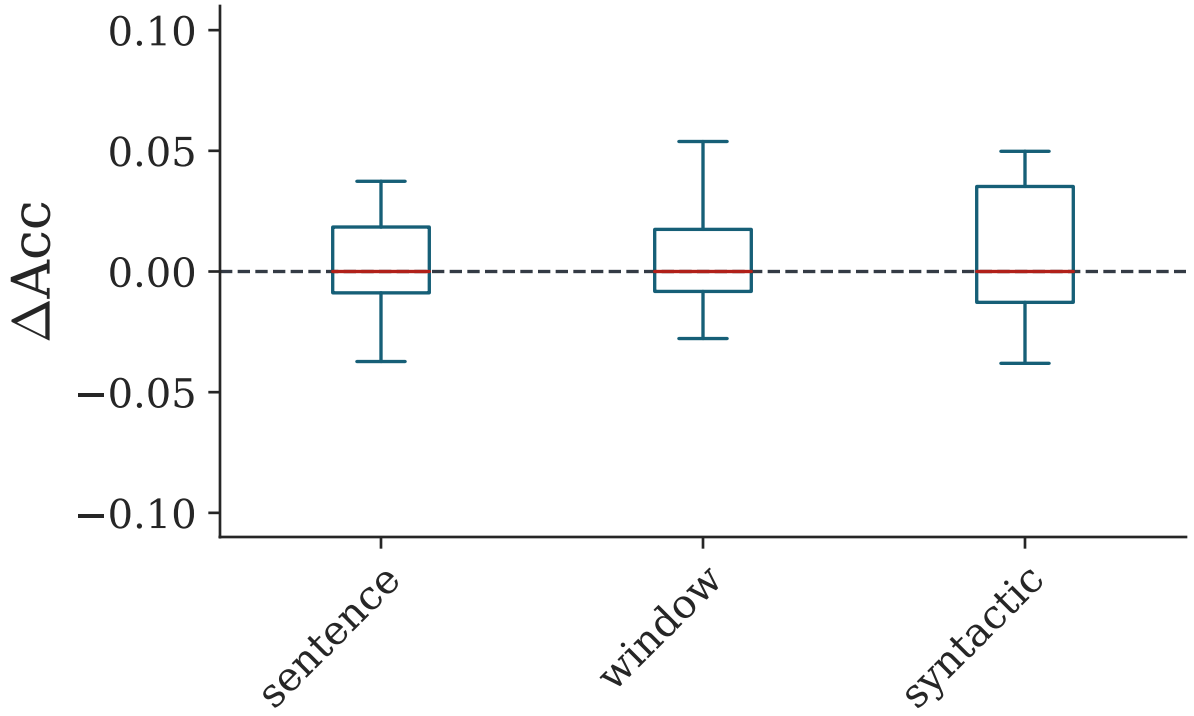

Figure 5: Distribution of multi-class accuracy differences between models trained with and without co-occurrences.

McRae (Spanish) — Multi-class

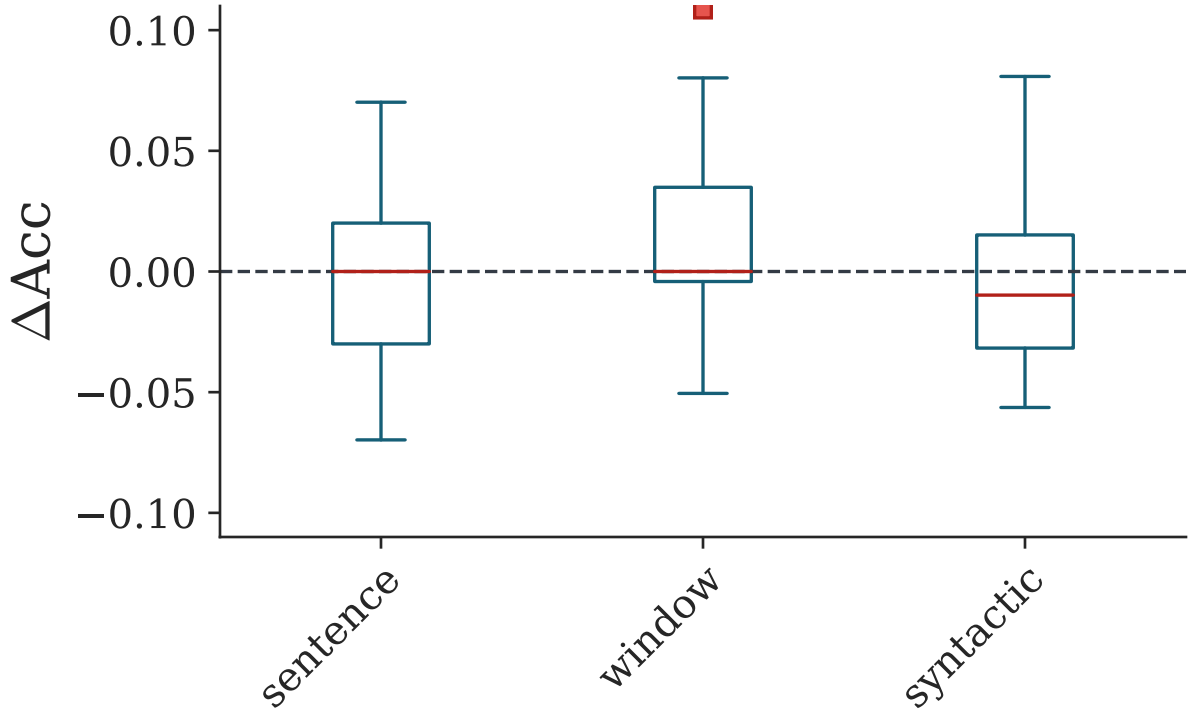

Figure 6: Distribution of multi-class accuracy differences between models trained with and without co-occurrences.

Binary Results

- Typically quite high (80-90% accuracy)
- Require negative examples (clearly not trivial)
- Similar to multi-class (no real distinct patterns regarding removal of co-occurrences).

Conclusions and so on

- No differences when removing co-occurrences
- Exception of colours for one dataset (English McRae)
- Cultivated a dataset for English and Spanish for different trait types

End

Thanks. Any questions, remarks, or criticisms?